

# Privacy Policies of Major LLM Providers

A Plain-Language Comparison: OpenAI (ChatGPT) · Google (Gemini) · Anthropic (Claude)

Compiled July 30, 2026 | Based on each company's publicly published privacy policies and help-center documentation

This report summarizes, in plain language, how three major AI providers describe their handling of user data across consumer and business/enterprise tiers: what is collected, whether conversations are used to train models, how long data is retained, and what controls are available to users. Policies change frequently — treat this as a general orientation and always check the provider's current policy page before making decisions involving sensitive, confidential, or regulated data.

## At a Glance

| Dimension                                            | OpenAI (ChatGPT)                                                                                                 | Google (Gemini)                                                                                         | Anthropic (Claude)                                                                                                             |
|------------------------------------------------------|------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------|
| Consumer default: used for training?                 | Yes, on Free/Plus, unless you opt out                                                                            | Yes, on the free consumer app, unless you turn off 'Gemini Apps Activity'                               | Yes, on Free/Pro/Max, unless you opt out (opt-in toggle introduced Aug–Sep 2025)                                               |
| Business/Enterprise default: used for training?      | No, not by default (API, Team, Enterprise)                                                                       | No, Workspace/Cloud enterprise data is excluded by default                                              | No, API, Claude for Work, and government tiers excluded by default                                                             |
| Consumer data retention (if you opt out of training) | Conversations kept until deleted; deleted chats fully removed within ~30 days                                    | Default 18 months (adjustable to 3 or 36 months) unless activity controls are changed                   | ~30 days on the back end after deletion, if training is off                                                                    |
| Consumer retention (if you opt in to training)       | Content may be retained per opted-in settings; export/delete tools available                                     | Governed by the same Gemini Apps Activity retention window you've set                                   | Up to 5 years for conversations used to improve models                                                                         |
| API / developer default                              | Not used for training; ~30-day abuse-monitoring retention; Zero Data Retention available for eligible enterprise | Prompts/responses logged briefly for safety and abuse detection; not used for training on paid tiers    | Not used for training; standard log retention 7 days (30 days optional via DPA); Zero Data Retention for qualifying enterprise |
| Human review possible?                               | Yes, on content not opted out, for safety and quality purposes                                                   | Only with organizational/user consent for Workspace; consumer app may be reviewed per activity settings | Yes, on opted-in content, and on safety-flagged or reported content regardless of setting                                      |
| Sells data to third parties?                         | No                                                                                                               | No                                                                                                      | No                                                                                                                             |
| Key compliance certifications (enterprise)           | SOC 2 Type 2; HIPAA BAA for eligible enterprise endpoints                                                        | ISO 27001/27017/27018; FedRAMP; HIPAA BAA for Workspace (with agreement)                                | SOC 2 Type II; HIPAA BAA available; FedRAMP in progress                                                                        |

Note: 'Default' describes the out-of-the-box setting for a typical new account as of mid-2026. Organizational admins, regional regulation (e.g., EU/UK GDPR), and specific contract terms can change these defaults.

## OpenAI — ChatGPT & the OpenAI API

OpenAI operates under a single Privacy Policy plus separate Usage Policies and an Enterprise Privacy page. The consumer product (ChatGPT Free and Plus) and the business/developer products (API, Team, Enterprise) follow opposite defaults for model training.

### Consumer tiers (Free, Plus)

- By default, conversations may be used to help improve and train OpenAI's models.
- Users can opt out under Settings → Data Controls → "Improve the model for everyone," or by using Temporary Chat, which is never saved to history or used for training.
- Once a conversation is deleted, it is removed from the visible history immediately and from OpenAI's backend systems within about 30 days, barring legal or security exceptions.

### Business tiers (Team, Enterprise, API)

- Content submitted through ChatGPT Team, ChatGPT Enterprise, and the standard API is not used to train OpenAI's models by default.
- API data retention defaults to a short window (about 30 days) for abuse monitoring, after which it is deleted.
- Zero Data Retention (ZDR) is available for qualifying enterprise customers, removing even that short-term storage.

### User controls & rights

OpenAI provides a self-service privacy portal for submitting a "do not train on my content" opt-out and data export requests, and a separate contact channel for formal data subject access requests (DSARs) under laws such as GDPR and CCPA.

### Compliance posture

OpenAI publishes SOC 2 Type 2 status and offers HIPAA Business Associate Agreements for eligible enterprise customers and endpoints, per its published Trust Portal.

---

## Google — Gemini App & Gemini for Workspace / Cloud

Google's approach is split across the consumer Gemini app (governed by the general Google Privacy Policy and "Gemini Apps Activity" settings) and Gemini for Workspace/Cloud/Vertex AI, which follows separate enterprise terms.

### Consumer Gemini app

- Conversations are generally stored under your Google Account and may be used to improve Google's products, including through human review, unless you disable Gemini Apps Activity.
- Default retention for saved activity is commonly cited around 18 months, adjustable to shorter (3 months) or longer (36 months) windows in account settings.
- Turning off Gemini Apps Activity limits both retention and use for improvement, but can also disable conversation history and related personalization features.

### Workspace, Cloud & Vertex AI (enterprise)

- Google states that enterprise prompts, files, and outputs within Gemini for Workspace and Gemini for Cloud are not used to train public models and are not reviewed by humans without permission.
- Workspace retention for Gemini prompts and responses is admin-configurable, commonly ranging from 90 days up to indefinite retention, with users able to delete individual conversations unless an admin restricts this.
- NotebookLM, a related Google AI product, is described as not retaining prompts/responses after a session ends.

### **User & admin controls**

Workspace administrators control retention windows, human-review permissions, and whether users can delete their own conversation history, via the Admin Console and Google Vault.

### **Compliance posture**

Google publishes ISO/IEC 27001, 27017, and 27018 certifications, FedRAMP authorization for government use, and offers a HIPAA Business Associate Addendum for Workspace with the appropriate agreement and admin configuration.

# Anthropic — Claude (Free/Pro/Max) & the Anthropic API / Claude for Work

Anthropic historically marketed itself as the privacy-first option, with no training on consumer conversations. In August–October 2025 it updated its Consumer Terms and Privacy Policy to introduce a model-training choice with a longer retention window for users who opt in.

## Consumer tiers (Free, Pro, Max)

- Since the 2025 policy update, new or resumed conversations may be used to help train future models if the user turns on the "help improve Claude" setting in Privacy Settings; this also applies to Claude Code sessions run from a consumer account.
- If a user opts in, Anthropic states it may retain that data in de-identified form for up to five years; deleting a conversation removes it from future training use.
- If a user does not opt in (or opts out), the prior behavior continues: roughly a 30-day backend retention window and no use in training.
- Anthropic states that safety-flagged or user-reported content may still be reviewed even when the general training setting is off.

## Business tiers (API, Claude for Work, Enterprise, government)

- Anthropic's Commercial Terms state that customer content submitted through the API, Claude for Work, Enterprise, and government offerings is not used to train models by default.
- Standard API log retention is commonly cited at about 7 days (reduced from a prior 30-day window), with an optional 30-day retention available through a Data Processing Addendum for organizations that need it for auditing.
- Zero Data Retention (ZDR) is available for qualifying enterprise customers.

## User controls & rights

Users manage the model-training toggle and review data settings under Settings → Privacy in Claude.ai, can delete individual conversations, and can request account-level data deletion or export in line with applicable law.

## Compliance posture

Anthropic states it does not sell user data to third parties, publishes SOC 2 Type II certification, offers a HIPAA Business Associate Agreement for qualifying enterprise deployments (used, for example, in some healthcare deployments), and reports FedRAMP authorization as in progress.

---

## Practical Takeaways

- **Free/consumer tiers of all three providers currently default to some form of data use for improving or training models**, and require the user to actively find and use an opt-out or activity-control setting.
- **Business, enterprise, and standard developer API usage across all three providers is excluded from model training by default** — this is the most consistent point of agreement between the three policies.
- **Retention windows vary widely**: short (days) for API traffic where zero data retention isn't used, months for consumer defaults, and up to several years for consumer accounts that opt in to training (notably Anthropic's 5-year window for opted-in data).
- **None of the three providers states that it sells user data to third parties.**

- **If handling client data, regulated data (health, financial, legal), or anything confidential, use the paid business/enterprise or API tier of any of these providers rather than the free consumer product, and confirm current terms directly with the vendor, since defaults and settings continue to change.**

---

Sources: OpenAI Privacy Policy, Usage Policies, and "How Your Data Is Used to Improve Model Performance" ([openai.com/policies](https://openai.com/policies)); Anthropic "Updates to Consumer Terms and Privacy Policy" and Commercial Terms ([anthropic.com/news](https://anthropic.com/news)); Google Workspace "Generative AI in Google Workspace Privacy Hub" and Google Cloud Gemini Enterprise documentation ([workspace.google.com](https://workspace.google.com), [cloud.google.com](https://cloud.google.com)). Independent policy trackers were also consulted for cross-checking; always verify against each provider's live policy page, as terms are revised periodically.

*This document is provided for general informational purposes only and is not legal advice. It is not a substitute for reading each provider's current, complete privacy policy and terms of service.*